

---

ID: 190041220

Tasfia Tasneem Annesha

CSE'19

## ▼ Importing Libraries

```
import numpy as np
from sklearn import datasets
from sklearn.model_selection import train_test_split
import matplotlib.pyplot as plt
import seaborn as sns
```

### ***class***

```
class LinearRegression:
    def __init__(self, lr = 0.001, n_iters = 100):
        self.lr = lr
        self.n_iters = n_iters
        self.weights = None
        self.bias = None

    def fit(self, X, y):
        # init parameters
        n_samples, n_features = X.shape
        self.weights = np.zeros(n_features)
        self.bias = 0
        self.mse_values = list([])

        for _ in range(self.n_iters):
            y_predicted = np.dot(X, self.weights) + self.bias
            self.mse_values.append(np.mean((y_predicted - y)**2))

            dw = (1/n_samples) * np.dot(X.T, (y_predicted - y))
            db = (1/n_samples) * np.sum(y_predicted - y)

            self.weights -= self.lr * dw
            self.bias -= self.lr * db

    def predict(self, X):
        y_predicted = np.dot(X, self.weights) + self.bias
        return y_predicted
```


```
X, y = datasets.make_regression(n_samples = 100, n_features = 1, noise = 20, random_state = 7)
print(X.shape,y.shape)
```

```
(100, 1) (100,)
```

y

```
array([ 7.74874748, 48.30195813, -99.86839046, 170.69731585,
 35.38423095, -24.26065523, 53.38022553, 40.93769612,
-58.37929858, -40.00080755, -5.72402446, -32.41843535,
-185.39492942, 16.10110541, 40.41032161, 65.57471889,
-91.91302503, 109.79950356, 20.67535726, -55.90394033,
-31.48725571, -34.18809808, -7.62204079, 0.66585166,
 6.58245201, -137.86783118, 28.96765275, -56.66561448,
-27.38686869, 2.47905123, 119.3324661 , -25.36827924,
 37.61826504, 15.71208367, 35.29998191, -47.8891707 ,
-115.4023037 , -6.29475666, -35.05780108, 135.50674081,
-43.00213356, 4.5838433 , -9.24111836, 90.55005602,
-42.13259317, 97.48433802, 138.54538857, 104.45930872,
 13.3261784 , -32.12098101, 51.52040062, -99.90017385,
 10.88356439, -22.16972415, 150.6949243 , 29.52128207,
-126.014095 , -5.93723276, -5.30531644, 71.23691077,
-42.3738208 , -46.50895613, -39.68083805, 123.18508718,
 27.1377505 , -26.09127408, -140.0584903 , -8.19702745,
-38.63819824, 76.11786241, 43.93083176, 31.19797233,
142.43536845, -112.89329896, -6.24848601, -0.64742201,
 89.92333305, 32.37393853, 54.52190665, -96.75561549,
-48.10701677, 48.43826367, 78.08997587, -104.44462323,
 8.99421674, 62.78311716, -149.07075314, -130.92984266,
-7.55831063, 99.78529963, -29.35974516, -55.73406339,
 12.65420352, -52.89627393, -127.55160577, 105.3439611 ,
-16.77527703, 44.58734207, 33.32282545, -8.41288496])
```

```
fig =plt.figure(figsize=(8,6))
plt.scatter(X[:,0],y,color='r',marker = 'o')
plt.show()
```



#test size 20%

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.2, random_state = 7)
```

```
X_train.shape
```

```
(80, 1)
```

```
regressor =LinearRegression(lr =0.001,n_iters =10000)
```

```
regressor.fit(X_train,y_train)
```

```
predicted= regressor.predict(X_test)
```

```
y_test
```

```
array([ -6.29475666,  28.96765275,  54.52190665, -55.73406339,
        -32.12098101,  65.57471889, -52.89627393,  31.19797233,
       -149.07075314,  -7.62204079,  16.10110541, -43.00213356,
         10.88356439, -185.39492942,  -7.55831063,  97.48433802,
       -32.41843535, -140.0584903 , -31.48725571,  20.67535726])
```

```
predicted
```

```
array([ -20.90190242,  41.38929807,  82.39416477, -72.82028561,
       -28.98147773,  46.65724654, -82.75515468,  22.80840514,
```

```
-124.93077671, -18.24788994, 36.8867518 , -51.54495826,  
-22.95720707, -158.5059681 , 17.31446106, 114.05935271,  
-32.81857505, -143.01289179, -20.9352896 , -9.32300296]]
```


```
def mse(y_true, y_predicted):  
    return np.mean((y_true - y_predicted)**2)
```

```
mse(predicted,y_test)
```

```
387.4268402120403
```


```
def draw_line (X_train, y_train, X_test, y_test, predicted):  
    fig =plt.figure(figsize=(8,6))  
    plt1 = plt.scatter(X_train, y_train, color = 'b')  
    plt2 = plt.scatter(X_test, y_test, color = 'r')  
    plt.plot(X_test,predicted,color='black',linewidth = 2)
```

```
draw_line(X_train, y_train, X_test, y_test, predicted)
```



```
plt.plot(regressor.mse_values)
```

```
[<matplotlib.lines.Line2D at 0x7fa281df0ca0>]
```



```
import pandas as pd
import seaborn as sns
!gdown 1m_uJkaKZvX24wdyAjXFIxTUY3SIawaCF
```

Downloading...

From: [https://drive.google.com/uc?id=1m\\_uJkaKZvX24wdyAjXFIxTUY3SIawaCF](https://drive.google.com/uc?id=1m_uJkaKZvX24wdyAjXFIxTUY3SIawaCF)

To: /content/CarPrice.csv

100% 26.7k/26.7k [00:00<00:00, 32.7MB/s]

## ▼ Load Prices in pandas data frame from a csv (excel) file


**Read data-** read\_csv()

In the following section the first 5 rows of the dataset are shown to give a brief overview of the

```
data = pd.read_csv('/content/CarPrice.csv')
data.head()
```

	car_ID	symboling	CarName	fueltype	aspiration	doornumber	carbody	dri
0	1	3	alfa-romero giulia	gas	std	two	convertible	
1	2	3	alfa-romero stelvio	gas	std	two	convertible	
2	3	1	alfa-romero Quadrifoglio	gas	std	two	hatchback	
3	4	2	audi 100 ls	gas	std	four	sedan	
4	5	2	audi 100ls	gas	std	four	sedan	

5 rows × 26 columns



wheelbase continuous

for some of the features list of possible values are fixed

shows the *columns and rows of the dataset*

```
data.shape
```

```
(205, 26)
```

## numericL features

```
data.describe()
```

	car_ID	symboling	wheelbase	carlength	carwidth	carheight	curbw
<b>count</b>	205.000000	205.000000	205.000000	205.000000	205.000000	205.000000	205.0
<b>mean</b>	103.000000	0.834146	98.756585	174.049268	65.907805	53.724878	2555.5
<b>std</b>	59.322565	1.245307	6.021776	12.337289	2.145204	2.443522	520.6
<b>min</b>	1.000000	-2.000000	86.600000	141.100000	60.300000	47.800000	1488.0
<b>25%</b>	52.000000	0.000000	94.500000	166.300000	64.100000	52.000000	2145.0
<b>50%</b>	103.000000	1.000000	97.000000	173.200000	65.500000	54.100000	2414.0
<b>75%</b>	154.000000	2.000000	102.400000	183.100000	66.900000	55.500000	2935.0
<b>max</b>	205.000000	3.000000	120.900000	208.100000	72.300000	59.800000	4066.0



Here we can see the names of the columns, their data types and the non-Null values count.

```
data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 205 entries, 0 to 204

Data columns (total 26 columns):

#	Column	Non-Null Count	Dtype
0	car_ID	205 non-null	int64
1	symboling	205 non-null	int64
2	CarName	205 non-null	object
3	fueltype	205 non-null	object
4	aspiration	205 non-null	object
5	doornumber	205 non-null	object
6	carbody	205 non-null	object
7	drivewheel	205 non-null	object
8	enginelocation	205 non-null	object
9	wheelbase	205 non-null	float64
10	carlength	205 non-null	float64
11	carwidth	205 non-null	float64
12	carheight	205 non-null	float64
13	curbweight	205 non-null	int64
14	enginetype	205 non-null	object
15	cylindernumber	205 non-null	object
16	enginesize	205 non-null	int64
17	fuelsystem	205 non-null	object
18	boreratio	205 non-null	float64
19	stroke	205 non-null	float64
20	compressionratio	205 non-null	float64
21	horsepower	205 non-null	int64
22	peakrpm	205 non-null	int64
23	citympg	205 non-null	int64
24	highwaympg	205 non-null	int64
25	price	205 non-null	float64

dtypes: float64(8), int64(8), object(10)

memory usage: 41.8+ KB

data.columns

```
Index(['car_ID', 'symboling', 'CarName', 'fueltype', 'aspiration',  
      'doornumber', 'carbody', 'drivewheel', 'enginelocation', 'wheelbase',  
      'carlength', 'carwidth', 'carheight', 'curbweight', 'enginetype',  
      'cylindernumber', 'enginesize', 'fuelsystem', 'boreratio', 'stroke',  
      'compressionratio', 'horsepower', 'peakrpm', 'citympg', 'highwaympg',  
      'price'],  
      dtype='object')
```



▼ **Data Cleaning:**

```
carnames = data['CarName']
CompanyNames = []

for carname in carnames:
    brand = carname.split(" ")[0]
    CompanyNames.append(brand)

data.drop(['CarName'], axis = 1, inplace = True)
data.insert(3, "CompanyName", CompanyNames)

data.head()
```

	car_ID	symboling	fueltype	CompanyName	aspiration	doornumber	carbody	dr
0	1	3	gas	alfa-romero	std	two	convertible	
1	2	3	gas	alfa-romero	std	two	convertible	
2	3	1	gas	alfa-romero	std	two	hatchback	
3	4	2	gas	audi	std	four	sedan	
4	5	2	gas	audi	std	four	sedan	

5 rows × 26 columns



```
data.CompanyName = data.CompanyName.str.lower()
```

```
data.CompanyName.unique()
```

```
array(['alfa-romero', 'audi', 'bmw', 'chevrolet', 'dodge', 'honda',
      'isuzu', 'jaguar', 'maxda', 'mazda', 'buick', 'mercury',
      'mitsubishi', 'nissan', 'peugeot', 'plymouth', 'porsche',
      'porcshe', 'renault', 'saab', 'subaru', 'toyota', 'toyouta',
      'vokswagen', 'volkswagen', 'vw', 'volvo'], dtype=object)
```

```
data.fueltype.unique()
```

```
array(['gas', 'diesel'], dtype=object)
```

***split fueltype coloumn in two different colums gas and diesel:***

## ▼ CREATE DUMMY VARIABLE COLUMNS

```
temp =pd.get_dummies(data['fueltype'])
```

```
temp.head()
```

	diesel	gas
0	0	1
1	0	1
2	0	1
3	0	1
4	0	1

## ▼ drop original fueltype data column

```
data.drop(['fueltype'],axis=1,inplace=True)
```

## ▼ CONCATINATE COLUMNS


axis=1 ->column

```
data=pd.concat([data,temp],axis=1)# as we are adding column we are concatinating coloum axis
```

```
data.head()
```

	car_ID	symboling	CompanyName	aspiration	doornumber	carbody	drivewheel
0	1	3	alfa-romero	std	two	convertible	rwd
1	2	3	alfa-romero	std	two	convertible	rwd
2	3	1	alfa-romero	std	two	hatchback	rwd
3	4	2	audi	std	four	sedan	fwd
4	5	2	audi	std	four	sedan	4wd

5 rows × 27 columns



## ▼ *Training model*

Training data set size 0.8 as here test set size is 0.2

```
data_train, data_test = train_test_split(data, test_size = 0.2)
```

```
len(X_train)
```

164

```
len(X_test)
```

41

```
X_train
```

```
array([[ 99.1, 186.6, 110. ,  0. ,  1. ],
       [104.5, 187.8, 156. ,  0. ,  1. ],
       [ 95.7, 169.7,  62. ,  0. ,  1. ],
       [ 93.7, 157.3,  68. ,  0. ,  1. ],
       [ 97.2, 173.4,  97. ,  0. ,  1. ],
       [101.2, 176.8, 101. ,  0. ,  1. ],
       [ 98.4, 176.2, 116. ,  0. ,  1. ],
       [113. , 199.6, 176. ,  0. ,  1. ],
       [101.2, 176.8, 121. ,  0. ,  1. ],
```



```
[102.4, 175.6, 73. , 1. , 0. ],
[ 93.7, 150. , 76. , 0. , 1. ],
[ 96.6, 180.3, 155. , 0. , 1. ],
[ 94.5, 155.9, 70. , 0. , 1. ],
[ 98.4, 175.7, 288. , 0. , 1. ],
[ 98.8, 177.8, 84. , 0. , 1. ],
[ 99.1, 186.6, 110. , 0. , 1. ],
[ 88.6, 168.8, 111. , 0. , 1. ],
[ 95.7, 158.7, 62. , 0. , 1. ],
[ 96.5, 175.4, 101. , 0. , 1. ],
[ 94.5, 168.7, 112. , 0. , 1. ],
[ 93.7, 150. , 76. , 0. , 1. ],
[ 93.7, 167.3, 68. , 0. , 1. ],
[ 95.9, 173.2, 145. , 0. , 1. ],
[ 96.3, 172.4, 116. , 0. , 1. ],
[106.7, 187.5, 123. , 1. , 0. ],
[ 93.1, 159.1, 68. , 0. , 1. ],
[102.4, 175.6, 92. , 0. , 1. ],
[102.7, 178.4, 175. , 0. , 1. ],
[ 98.4, 176.2, 116. , 0. , 1. ],
[105.8, 192.7, 140. , 0. , 1. ],
[102.4, 175.6, 92. , 0. , 1. ],
[ 93.7, 157.3, 68. , 0. , 1. ],
[107.9, 186.7, 97. , 0. , 1. ],
[102.4, 175.6, 92. , 0. , 1. ],
[ 93.7, 157.9, 73. , 0. , 1. ],
[114.2, 198.9, 97. , 0. , 1. ],
[ 94.5, 165.3, 69. , 0. , 1. ],
[ 95.3, 169. , 101. , 0. , 1. ],
[ 95.7, 158.7, 62. , 0. , 1. ],
[112. , 199.2, 184. , 0. , 1. ],
[115.6, 202.6, 155. , 0. , 1. ],
[ 93.7, 157.3, 68. , 0. , 1. ],
[ 93.1, 166.8, 68. , 0. , 1. ],
[ 99.8, 177.3, 110. , 0. , 1. ],
[ 98.4, 176.2, 116. , 0. , 1. ],
[ 94.5, 165.3, 69. , 0. , 1. ],
[109.1, 188.8, 106. , 1. , 0. ],
[ 98.8, 177.8, 84. , 0. , 1. ],
[ 98.4, 176.2, 116. , 0. , 1. ],
[ 95.7, 166.3, 70. , 0. , 1. ],
[109.1, 188.8, 134. , 0. , 1. ],
[102.4, 175.6, 92. , 0. , 1. ],
[ 94.5, 168.7, 70. , 0. , 1. ],
[114.2, 198.9, 95. , 1. , 0. ],
[ 94.5, 168.7, 112. , 0. , 1. ],
[ 94.5, 165.6, 69. , 0. , 1. ],
[ 97.3, 171.7, 100. , 0. , 1. ],
[ 96.3, 172.4, 88. , 0. , 1. ],
```

```
data_train.head()
```

	car_ID	symboling	CompanyName	aspiration	doornumber	carbody	drivewheel
<b>132</b>	133	3	saab	std	two	hatchback	fwd
<b>181</b>	182	-1	toyouta	std	four	wagon	rwd
<b>155</b>	156	0	toyota	std	four	wagon	4wd
<b>25</b>	26	1	dodge	std	four	sedan	fwd
<b>99</b>	100	0	nissan	std	four	hatchback	fwd

5 rows × 27 columns

```
feature_columns = ['wheelbase', 'carlength', 'horsepower', 'diesel', 'gas']
X_train = data_train[feature_columns].values
X_test = data_test[feature_columns].values
y_train = data_train['price'].values
y_test = data_test['price'].values
X_train.shape
```

(164, 5)

## ▼ creating linear regression object-'regressor'

object creation for linear regression from sklearn linear\_model

then fitting data->training model

1st arg- 2d array

2nd arg- y axis

then predicting prices

```
regressor = LinearRegression(lr = 0.000000000000001, n_iters = 1000)
regressor.fit(X_train, y_train)
predicted = regressor.predict(X_test)
```

```
regressor.predict(330)
```

```
array([4.45858266e-04, 7.90347644e-04, 5.40027085e-04, 5.80231247e-07,
       3.85184042e-06])
```

```
mse(predicted y test)
```

```
mse(predicted, y_test)
```

```
222716164.80540848
```

[Colab paid products](#) - [Cancel contracts here](#)

---

✓ 0s completed at 12:58 AM

